

A short talk about efficient C++ programming

*„C++11 feels like a new language:
The pieces just fit together better than they used to
and I find a higher-level style of programming
more natural than before and as efficient as ever.”*

Bjarne Stroustrup

*Mateusz Pusz
Intel Technology Poland*




To the extent possible under law, Intel Corporation has waived all copyright and related or neighboring rights to this work.

- * Some images used are the subject of other licensing terms thus are not released to the public domain.

Plan for next 45 minutes

- **We will scope on efficient programming in C++ with special attention on C++11**
 - Can C++ be efficient at all???
- **I will not bore you with C++ reference documentation for any specific C++11 feature**
 - There are already good books, blogs and Wikipedia out there
- **However I will showcase you a lot of new C++11 features**
- **I want to show you how different aspects of C++ computer programming interact with each other**
 - C++ is a programming language and not just a collection of separate features
- **Hopefully we will have some fun too!**
- **And ...**

*



... it is never too late
to play with yet
another algorithm 😊

Simple programming problem

'my_int' is a simple class containing 'int' as a member and a few user defined methods. 'int' member inside 'my_int' class should be initialized during class construction.

Implement a function 'test' that:

- *takes 2 integral arguments 'n' and 'k'*
- *creates a container of 'n' 'my_int' objects with random values assigned*
- *sorts created container*
- *inserts a new random value keeping the container sorted (repeat 'k' times)*
- *verifies if the resulting container is sorted*

'my_int' is a simple class containing 'int' as a member and a few user defined methods

```
class my_int {  
    int _value;  
public:  
  
    // some methods...  
};
```

'int' member inside 'my_int' class should be initialized during class construction

```
class my_int {  
    int _value;  
public:  
    explicit constexpr my_int(int value): _value{value} {}  
    // some methods...  
};
```

generalized constant expression

uniform initialization

implement a function 'test' that ...

```
class my_int {
    int _value;
public:
    explicit constexpr my_int(int value): _value{value} {}
    // some methods...
};

void test(int n, int k)
{
    // create container of 'n' random values
    // sort container
    // k times insert a new random value keeping the container sorted
    // verify if sorted
}
```

“sorting”, “comparing values”, ... → I need to extend ‘my_int’ interface

```
class my_int {  
    int _value;  
public:  
    explicit constexpr my_int(int value): _value{value} {}  
    // some methods...  
  
    // returns stored value  
    explicit constexpr operator int() const { return _value; }  
  
    // user-defined less-comparator needed for sorting  
    friend constexpr bool operator <(const my_int& left,  
                                    const my_int& right)  
    { return left._value < right._value; }  
};
```



“class”, not a fundamental type → operator new

“inserting in the middle”, “sorting”, ... → std::list

```
using elem_type = my_int*;  
using container = std::list<elem_type>;
```



type aliases

“random values” → provide random numbers generator

```
int generate()
{
    ... // not enough time today to describe the implementation
}
```

XKCD way ...

```
int getRandomNumber()
{
    return 4; // chosen by fair dice roll.
              // guaranteed to be random.
}
```

*

... or use C++11 <random>

creates a container of 'n' 'my_int' objects with random values assigned

```
int generate();          // implementation provided in Backup
```

```
void init_random(int num, container& values)
{
    for(int i=0; i<num; ++i)
        values.emplace_back(new my_int{generate()});
}
```

faster container operations based on Perfect Forwarding

```
void test(int n, int k)
{
    // create container of 'n' random values
    container values;
    init_random(n, values);
    // ...
}
```

$O(n)$

sorts created container

```
void test(int n, int k)
{
    // ...
    // sort container
    auto ptr_cmp = [](const elem_type& l, const elem_type& r)
                    { return *l < *r; };
    values.sort(ptr_cmp);
    // ...
}
```

type inference

lambda -> [captures](args){ body; }


$O(n \log n)$

inserts a new random value keeping the container sorted (repeat 'k' times)

```
void test(int n, int k)
{
    // ...

    // k times insert a new random value keeping the container sorted
    for(int i=0; i<k; ++i) {
        const int val = generate();
        values.emplace(find_if(begin(values), end(values),
                                [&](const elem_type& v)
                                { return static_cast<int>(*v) > val; })),
                        new my_int{val});
    }

    // ...
}
```



faster container operations based on Perfect Forwarding

$O(k * (n+k))$

verifies if the resulting container is sorted

```
void test(int n, int k)
{
    // ...
    auto ptr_cmp = [](const elem_type& l, const elem_type& r)
        { return *l < *r; };
    // ...

    // verify if sorted
    std::cout << (is_sorted(begin(values), end(values), ptr_cmp) ?
        "PASS\n" : "FAIL\n");
}
```

new C++11 algorithm

$O(n)$

```
$ ./test
PASS
```

Are we there yet?

Whole code

Do you see a problem?

```
void test(int n, int k)
{
    // create container of 'n' random values
    container values;
    init_random(n, values);

    // sort container
    auto ptr_cmp = [](const elem_type& l, const elem_type& r){ return *l < *r; };
    values.sort(ptr_cmp);

    // k times insert a new random value keeping the container sorted
    for(int i=0; i<k; ++i) {
        const int val = generate();
        values.emplace(find_if(begin(values), end(values),
                                [&](const elem_type& v)
                                { return static_cast<int>(*v) > val; })),
                        new my_int{val});
    }

    // verify if sorted
    std::cout << (is_sorted(begin(values), end(values), ptr_cmp) ? "PASS\n":"FAIL\n");
}
```

Whole code

Do you see a problem?

```
void test(int n, int k)
{
    // create container of 'n' random values
    container values;
    init_random(n, values);

    // sort container
    auto ptr_cmp = [](const elem_type& l, const elem_type& r){ return *l < *r; };
    values.sort(ptr_cmp);

    // k times insert a new random value keeping the container sorted
    for(int i=0; i<k; ++i) {
        const int val = generate();
        values.emplace(find_if(begin(values), end(values),
                                [&](const elem_type& v)
                                { return static_cast<int>(*v) > val; })),
                        new my_int{val});
    }

    // verify if sorted
    std::cout << (is_sorted(begin(values), end(values), ptr_cmp) ? "PASS\n":"FAIL\n");
}
```



**We have a memory
leak here**

Adding cleanup is simple, right?

```
void test(int n, int k)
{
    // ...

    // cleanup the container
    for(auto& elem : values)
        delete elem;
}
```



range-based for loop

Whole code

Do you see a problem?

```
void test(int n, int k)
{
    // create container of 'n' random values
    container values;
    init_random(n, values);

    // sort container
    auto ptr_cmp = [](const elem_type& l, const elem_type& r){ return *l < *r; };
    values.sort(ptr_cmp);

    // k times insert a new random value keeping the container sorted
    for(int i=0; i<k; ++i) {
        const int val = generate();
        values.emplace(find_if(begin(values), end(values),
                               [&](const elem_type& v)
                               { return static_cast<int>(*v) > val; })),
                       new my_int{val});
    }

    // verify if sorted
    std::cout << (is_sorted(begin(values), end(values), ptr_cmp) ? "PASS\n" : "FAIL\n");

    // cleanup the container
    for(auto& elem : values)
        delete elem;
}
```

Whole code

Do you see a problem?

```
void test(int n, int k)
{
    // create container of 'n' random values
    container values;
    init_random(n, values);

    // sort container
    auto ptr_cmp = [](const elem_type& l, const elem_type& r){ return *l < *r; };
    values.sort(ptr_cmp);

    // k times insert a new random value keeping the container sorted
    for(int i=0; i<k; ++i) {
        const int val = generate();
        values.emplace(find_if(begin(values), end(values),
                                [&](const elem_type& v)
                                { return static_cast<int>(*v) > val; })),
                        new my_int{val});
    }

    // verify if sorted
    std::cout << (is_sorted(begin(values), end(values), ptr_cmp) ? "PASS\n" : "FAIL\n");

    // cleanup the container
    for(auto& elem : values)
        delete elem;
}
```



Resource Acquisition Is Initialization (RAII)

```
class MyResource {  
    // private resource  
public:  
    MyResource(/* args */)   
    { // do whatever is needed to obtain ownership over resource }  
    ~MyResource()  
    { // do whatever is needed to reclaim the resource }  
    // more class stuff here  
};
```

*If you dissect the words of the **RAII** acronym (Resource Acquisition Is Initialization), you will think RAII is about acquiring resources during initialization. However the power of RAII comes not from tying **acquisition** to **initialization**, but from tying **reclamation** to **destruction**.*

Bjarne Stroustrup

RAII usage in C++ Standard Library

■ Smart Pointers

- `std::unique_ptr<T, Deleter = std::default_delete<T>>`
- `std::shared_ptr<T>`
- `std::weak_ptr<T>`
- ~~`std::auto_ptr<T>`~~

■ Containers

- all STL containers manage ownership of their data
- `std::string` also became a regular container in C++11

■ File streams

■ Mutex locks

- `std::lock_guard<Mutex>`
- `std::unique_lock<Mutex>`

■ More...

Let's fix our code to be C++ exception safe

```
using elem_type = std::unique_ptr<my_int>;  
// ...
```

```
void test(int n, int k)  
{  
    // ...
```

```
// cleanup the container  
for(auto &elem : values)  
    delete elem;  
}
```

**Writing exception safe code is not
that hard at all!!!**

Tweaking 'n' in 'f(n)'

Can we make our code faster?

Solution	test(100 000, 10 000)	test(10 000 000, 0)	test(0, 100 000)
Original	25,5s	14,3s	63,2s
Change 1	9,75s	9,6s	30,4s
Change 2	0,52s	0,88s	2,4
Change 3	0,33s	0,88s	1,5s
Speedup	77x	16x	42x

Yes, A LOT faster

“class”, not a fundamental type → operator new
“inserting in the middle”, “sorting”, ... → std::list

```
using elem_type = my_int*;  
using container = std::list<elem_type>;
```

WRONG!!!

C++ is not C# or Java

■ Heap usage should be avoided if possible

- allocation and deallocation of heap memory is slow
- obtaining data from non-cached memory is slow
- heap allocation can fail
- allocation of many small objects causes huge memory fragmentation

■ C++ loves value semantics

- using pointers changes semantics of copy, assignment and equality
- pointer dereferencing takes time
- much easier to write thread-safe code
- reference/pointer semantics causes aliasing problems
 - *compiler optimizer cannot do its best -> slower code*
- Copy Elision (RVO) and Move Semantics improve things a lot

Change 1

Avoid pointers

```
using elem_type = my_int;
// ...

void test(int n, int k)
{
    // ...
    auto ptr_cmp = [](const elem_type& l, const elem_type& r)
        { return *l < *r; };
    // ...
}
```

1.5x-2x speedup

Change 2

Use `std::vector<T>` as a default container

```
using container=std::vector<elem_type>;

container init_random(int n, int k)
{
    container values;
    values.reserve(n + k);
    for(int i=0; i<n; ++i)
        values.emplace_back(generate());
    return values;
}
```

```
void test(int n, int k)
{
    // create container of 'n' random values
    auto values = init_random(n, k);

    // sort container
    sort(begin(values), end(values));

    // ...
}
```


10x-18x speedup. Why?

Algorithmic point of view

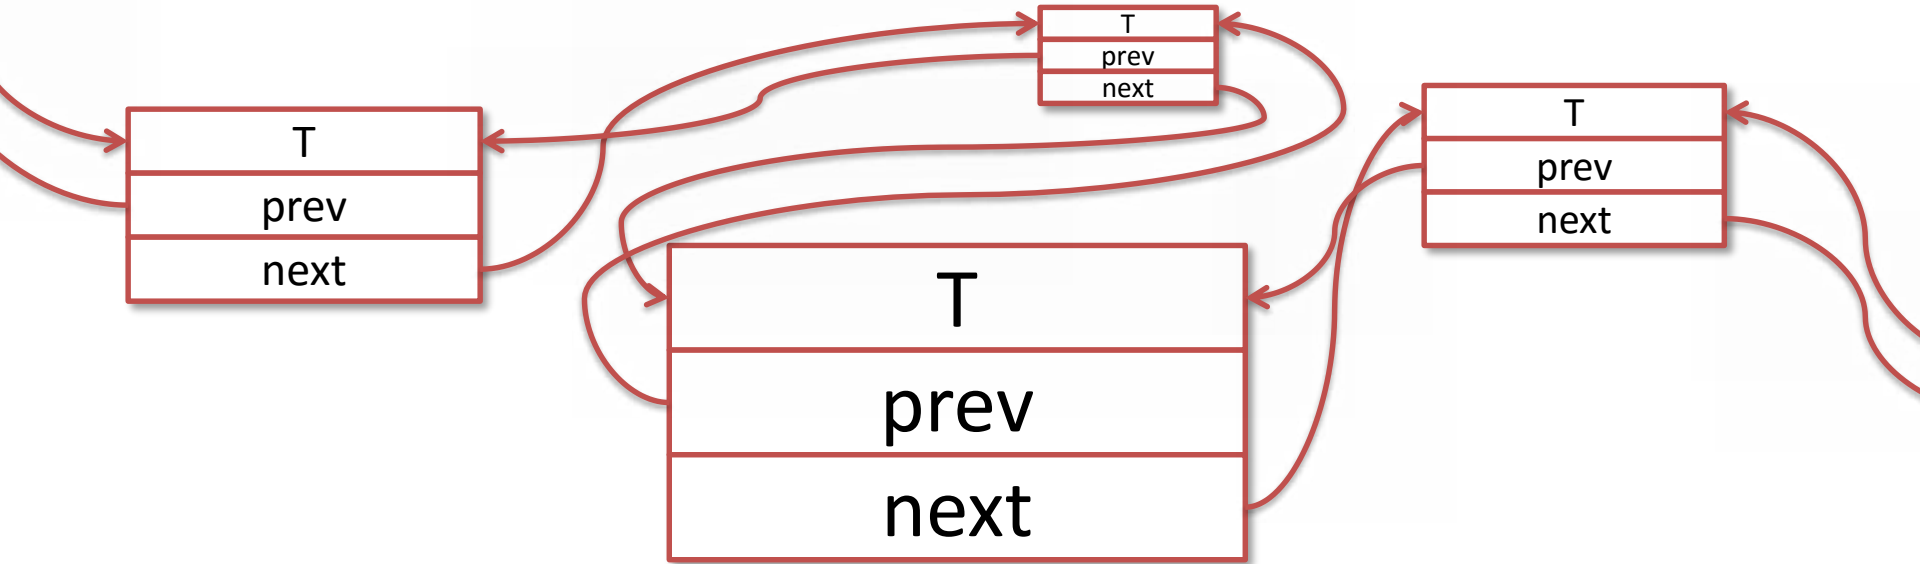
Operation	<code>std::list<T></code>	<code>std::vector<T></code>
create container of 'n' random values	$O(n)$	$O(n)$
sort container	$O(n \log n)$	$O(n \log n)$
'k' times insert a new random value keeping the container sorted	$O(k * (n+k))$	$O(k * (n+k))$
verify if sorted	$O(n+k)$	$O(n+k)$
Overall complexity	$O(k * (n+k))$	$O(k * (n+k))$

10x-18x speedup. Why?

std::list<T>




std::list<T>



- easy to insert new elements
- easy to erase existing elements
- easy to reorder elements
- bidirectional iteration
- memory ineffective
 - for small `T` large memory overhead
 - a lot of dynamic memory allocations and deallocations
 - a lot of pointer dereferences during iteration
 - not cache friendly


std::list<int> iteration

CPU cache



M
E
M
O
R
Y


std::vector<T>



- **inserting new elements**
 - end – easy (unless buffer reallocation is needed)
 - begin, middle – hard
- **erasing existing elements**
 - end – easy
 - begin, middle – hard
- **swapping values needed to reorder elements**
- **random access iteration**
- **memory effective**
 - small memory overhead (reserved space)
 - nearly no dynamic memory allocations and deallocations
 - no pointer dereferences during iteration
 - cache friendly!!!

std::vector<int> iteration

CPU cache



Writing cache friendly
code makes a huge
difference!

M
E
M
O
R
Y

A new possibility for algorithmic optimization

```
void test(int n, int k)
{
    // create container of 'n' random values
    // sort container
    // k times insert a new random value keeping the container sorted
    // verify if sorted
}
```

§25.4.3.1

template<class ForwardIterator, class T>

*ForwardIterator **lower_bound**(ForwardIterator **first**, ForwardIterator **last**, const T& **value**);*

1 **Requires:** The elements e of $[first, last)$ shall be partitioned with respect to the expression $e < value$

...

3 **Complexity:** At most $\log_2(last - first) + O(1)$ comparisons.

Change 3

Know and use C++ algorithms

```
void test(int n, int k)
{
    // ...
    // k times insert a new random value keeping the container sorted
    for(int i=0; i<k; ++i) {
        const my_int val{generate()};
        values.emplace(lower_bound(begin(values), end(values), val),
                        val);
    }
    // ...
}
```

1x-1.6x speedup

Whole code

```
std::vector<my_int> init_random(int n, int k)
{
    std::vector<my_int> values;
    values.reserve(n + k);
    for(int i=0; i<n; ++i)
        values.emplace_back(generate());
    return values;
}
```

```
void test(int n, int k)
{
    auto values = init_random(n, k);
    sort(begin(values), end(values));

    // k times insert a new random value keeping the container sorted
    for(int i=0; i<k; ++i) {
        const my_int val{generate()};
        values.emplace(lower_bound(begin(values), end(values), val), val);
    }

    std::cout << (is_sorted(begin(values), end(values)) ? "PASS\n" : "FAIL\n");
}
```


Isn't it a great example of efficient programming?

TAKE AWAYS



Learn C++ to be a more efficient programmer

- containers,
- algorithms,
- new C++11/C++14 features



Use RAII design pattern as your first resort tool for management of all resources



Use tools provided
by the C++ standard.

Do not reinvent the
wheel!!!



Remember that
memory access is the
bottleneck of many
today's applications

- limit pointers usage
- write cache friendly code



Read about
Copy Elision,
Move Semantics,
Perfect Forwarding,
Constant Expressions
to write faster
C++ code





Questions?



Thank you

Happy coding!!!

Backup

```

std::function<int()> make_random_int_generator(
    int min = std::numeric_limits<int>::min(),
    int max = std::numeric_limits<int>::max(),
    std::mt19937::result_type seed = std::mt19937::default_seed)
{
    std::mt19937 gen{seed};
    std::uniform_int_distribution<int> distr{min, max};
    return [=]() mutable { return distr(gen); };
}

int generate()
{
    static auto generator = make_random_int_generator();
    return generator();
}

```