

Small Lie in Big O

Truth

Mateusz Pusz

November 15, 2016

SIMPLE PROGRAMMING PROBLEM

'sample' is a simple class containing 8-bit unsigned integral member and a few user defined methods. Integral member inside 'sample' class should be initialized during class construction.

Implement function 'test' that:

- takes 2 integral arguments 'n' and 'k'*
- creates a collection of 'n' 'sample' objects with random values assigned*
- sorts created collection*
- inserts a new random value keeping the collection sorted (repeat 'k' times)*
- returns average value from all the samples in the collection*

'sample' is a simple class containing 8-bit unsigned integral member and a few user defined methods

```
class sample {  
    std::uint8_t value_;  
public:  
  
    // more interface...  
};
```

Integral member inside 'sample' class should be initialized during class construction

```
class sample {  
    std::uint8_t value_;  
public:  
    explicit constexpr sample(std::uint8_t value) noexcept : value_{value} {}  
    // more interface...  
};
```

implement a function 'test' that takes 2 integral arguments 'n' and 'k'

```
class sample {  
    std::uint8_t value_;  
public:  
    explicit constexpr sample(std::uint8_t value) noexcept : value_{value} {}  
    // more interface...  
};  
  
std::uint8_t test(int n, int k)  
{  
  
}
```

implement a function 'test' that ...

```
class sample {
    std::uint8_t value_;
public:
    explicit constexpr sample(std::uint8_t value) noexcept : value_{value} {}
    // more interface...
};

std::uint8_t test(int n, int k)
{
    // create collection of 'n' random samples
    // sort the collection
    // k times insert a new random sample keeping the collection sorted
    // return average value from samples in the collection
}
```

“sorting”, “comparing values”, ... -> need to extend ‘sample’ interface

```
class sample {
    std::uint8_t value_;
public:
    explicit constexpr sample(std::uint8_t value) noexcept : value_{value} {}
    // more interface...

    // returns stored value
    constexpr operator std::uint8_t() const noexcept { return value_; }
};
```


“class”, not a fundamental type -> dynamic memory allocation

“inserting in the middle”, “sorting”, ... -> std::list

```
using elem_type = sample*;  
using collection = std::list<elem_type>;
```


“random values” -> provide random numbers generator

```
std::uint8_t generate()  
{  
    ... // probably not enough time today to describe the implementation  
}
```

XKCD way ...

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

... or use C++11 `<random>`

creates a collection of 'n' 'sample' objects with random values assigned

```
collection init_random(int n)
{
    collection samples;
    for(int i=0; i<num; ++i)
        samples.emplace_back(new sample{generate()});
    return samples;
}

std::uint8_t test(int n, int k)
{
    // create collection of 'n' random samples
    auto samples = init_random(n);

    // ...
}
```

$O(n)$

sorts created collection

```
std::uint8_t test(int n, int k)
{
    // create collection of 'n' random samples
    auto samples = init_random(n);

    // sort the collection
    values.sort([](const auto& l, const auto& r) { return *l < *r; });

    // ...
}
```

$O(n \log n)$

inserts a new random value keeping the collection sorted (repeat 'k' times)

```
std::uint8_t test(int n, int k)
{
    // ...

    // k times insert a new random sample keeping the collection sorted
    for(int i=0; i<k; ++i) {
        const sample new_sample{generate()};
        samples.emplace(find_if(begin(samples), end(samples),
                                [&](const auto& s) { return new_sample < *s; })),
                        new sample{new_sample});
    }

    // ...
}
```

$O(k * (n+k))$

returns average value from all the samples in the collection

```
std::uint8_t test(int n, int k)
{
    // ...

    // return average value from samples in the collection
    return std::accumulate(begin(samples), end(samples), std::uint64_t{},
        [](const auto& sum, const auto& s){ return sum + *s; }) / samples.size();
}
```

$O(n)$

Whole code

Do you see a problem?

```
std::uint8_t test(int n, int k)
{
    // create collection of 'n' random samples
    auto samples = init_random(n);

    // sort the collection
    samples.sort([](const auto& l, const auto& r) { return *l < *r; });

    // k times insert a new random sample keeping the collection sorted
    for(int i = 0; i < k; ++i) {
        const sample new_sample{generate()};
        samples.emplace(find_if(begin(samples), end(samples),
                                [&](const auto& s) { return new_sample < *s; })),
                        new sample{new_sample});
    }

    // return average value from samples in the collection
    return std::accumulate(begin(samples), end(samples), std::uint64_t{},
                            [](const auto& sum, const auto& s){ return sum + *s; }) / samples.size();
}
```

Whole code

Do you see a problem?

```
std::uint8_t test(int n, int k)
{
    // create collection of 'n' random samples
    auto samples = init_random(n);

    // sort the collection
    samples.sort([](const auto& l, const auto& r) { return *l < *r; });

    // k times insert a new random sample keeping the collection sorted
    for(int i = 0; i < k; ++i) {
        const sample new_sample{generate()};
        samples.emplace(find_if(begin(samples), end(samples),
                                [&](const auto& s) { return new_sample < *s; })),
                        new sample{new_sample});
    }

    // return average value from samples in the collection
    return std::accumulate(begin(samples), end(samples), std::uint64_t{},
                            [](const auto& sum, const auto& s){ return sum + *s; }) / samples.size();
}
```




**We have a memory
leak here**

Adding cleanup is simple, right?

```
std::uint8_t test(int n, int k)
{
    // ...

    // cleanup the collection
    for(auto& s : samples)
        delete s;
}
```

Whole code

Do you see a problem?

```
std::uint8_t test(int n, int k)
{
    // create collection of 'n' random samples
    auto samples = init_random(n);

    // sort the collection
    samples.sort([](const auto& l, const auto& r) { return *l < *r; });

    // k times insert a new random sample keeping the collection sorted
    for(int i = 0; i < k; ++i) {
        const sample new_sample{generate()};
        samples.emplace(find_if(begin(samples), end(samples),
                                [&](const auto& s) { return new_sample < *s; }),
                        new sample{new_sample});
    }

    // return average value from samples in the collection
    auto avg = std::accumulate(begin(samples), end(samples), std::uint64_t{},
                               [](const auto& sum, const auto& s){ return sum + *s; }) / samples.size();

    // cleanup the collection
    for(auto& s : samples)
        delete s;

    return avg;
}
```

Whole code

Do you see a problem?

```
std::uint8_t test(int n, int k)
{
    // create collection of 'n' random samples
    auto samples = init_random(n);

    // sort the collection
    samples.sort([](const auto& l, const auto& r) { return *l < *r; });

    // k times insert a new random sample keeping the collection sorted
    for(int i = 0; i < k; ++i) {
        const sample new_sample{generate()};
        samples.emplace(find_if(begin(samples), end(samples),
                                [&](const auto& s) { return new_sample < *s; }),
                        new sample{new_sample});
    }

    // return average value from samples in the collection
    auto avg = std::accumulate(begin(samples), end(samples), std::uint64_t{},
                                [](const auto& sum, const auto& s){ return sum + *s; }) / samples.size();

    // cleanup the collection
    for(auto& s : samples)
        delete s;

    return avg;
}
```



The code is not exception safe!!!

RESOURCE ACQUISITION IS INITIALIZATION (RAII)

```
class resource {  
    // resource handle  
public:  
    resource(/* args */)   
    { /* obtain ownership of a resource and store the handle */ }  
    ~resource()  
    { /* reclaim the resource */ }  
};
```

If you dissect the words of the RAII acronym (Resource Acquisition Is Initialization), you will think RAII is about acquiring resources during initialization. However the power of RAII comes not from tying acquisition to initialization, but from tying reclamation to destruction.

-- Bjarne Stroustrup

RAII USAGE IN C++ STANDARD LIBRARY

- Smart Pointers
 - `std::unique_ptr<T, Deleter = std::default_delete<T>>`
 - `std::shared_ptr<T>`
 - `std::weak_ptr<T>`
 - ~~`std::auto_ptr<T>`~~
- Containers
 - all STL containers manage ownership of their data
- File streams
- Mutex locks
 - `std::lock_guard<Mutex>`
 - `std::unique_lock<Mutex>`
- More...

Let's fix our code to be C++ exception safe

```
using elem_type = std::unique_ptr<sample>;  
// ...  
  
// replace all  
new sample{...}  
// with  
std::make_unique<sample>(...)  
  
std::uint8_t test(int n, int k)  
{  
    // ...  
  
    // cleanup the collection  
    for(auto& s : samples)  
        delete s;  
}
```

Writing exception safe code
is not that hard at all!!!

TWEAKING 'N' IN 'F(N)'

Can we make our code faster?

Solution	test(10 000 000, 0)	test(0, 100 000)	test(100 000, 10 000)
Original	10,7s	104,5s	32,6s

TWEAKING 'N' IN 'F(N)'

Can we make our code faster?

Solution	test(10 000 000, 0)	test(0, 100 000)	test(100 000, 10 000)
Original	10,7s	104,5s	32,6s
Change #1	9,3s	57,1s	16,4s
Change #2	0,75s	1,7s	0,37s
Change #3	0,75s	0,98s	0,21s
Speedup	14x	106x	155x

Yes, A LOT faster

“class”, not a fundamental type -> dynamic memory allocation
“inserting in the middle”, “sorting”, ... -> std::list

```
using elem_type = sample*;  
using collection = std::list<elem_type>;
```

WRONG!!!

C++ IS NOT C# OR JAVA

- Heap usage should be avoided if possible
 - allocation and deallocation of heap memory is slow
 - obtaining data from non-cached memory is slow
 - heap allocation can fail
 - allocation of many small objects causes huge memory fragmentation
- C++ loves value semantics
 - using pointers changes semantics of copy, assignment and equality
 - pointer dereferencing takes time
 - much easier to write thread-safe code
 - reference/pointer semantics causes aliasing problems
 - compiler optimizer cannot do its best -> slower code
 - Copy Elision (RVO) and Move Semantics improve things a lot

CHANGE #1

Avoid pointers

```
using elem_type = sample;
// ...

std::uint8_t test(int n, int k)
{
    // ...
    samples.sort([](const auto& l, const auto& r) { return *l < *r; });
    // ...
    auto avg = std::accumulate(begin(samples), end(samples), std::uint64_t{},
                               [](const auto& sum, const auto& s){ return sum + *s; }) / samples.size();
    // ...
}
```

1.15x-2x speedup

CHANGE #2

Use `std::vector<T>` as a default container

```
using collection = std::vector<elem_type>;

collection init_random(int n, int k)
{
    collection samples;
    samples.reserve(n + k);
    for(int i=0; i<n; ++i)
        samples.emplace_back(generate());
    return samples;
}
```

```
std::uint8_t test(int n, int k)
{
    // create collection of 'n' random samples
    auto samples = init_random(n, k);
    // ...

    // sort the collection
    sort(begin(samples), end(samples));
    // ...
}
```


12x-44x speedup. Why?

ALGORITHMIC POINT OF VIEW

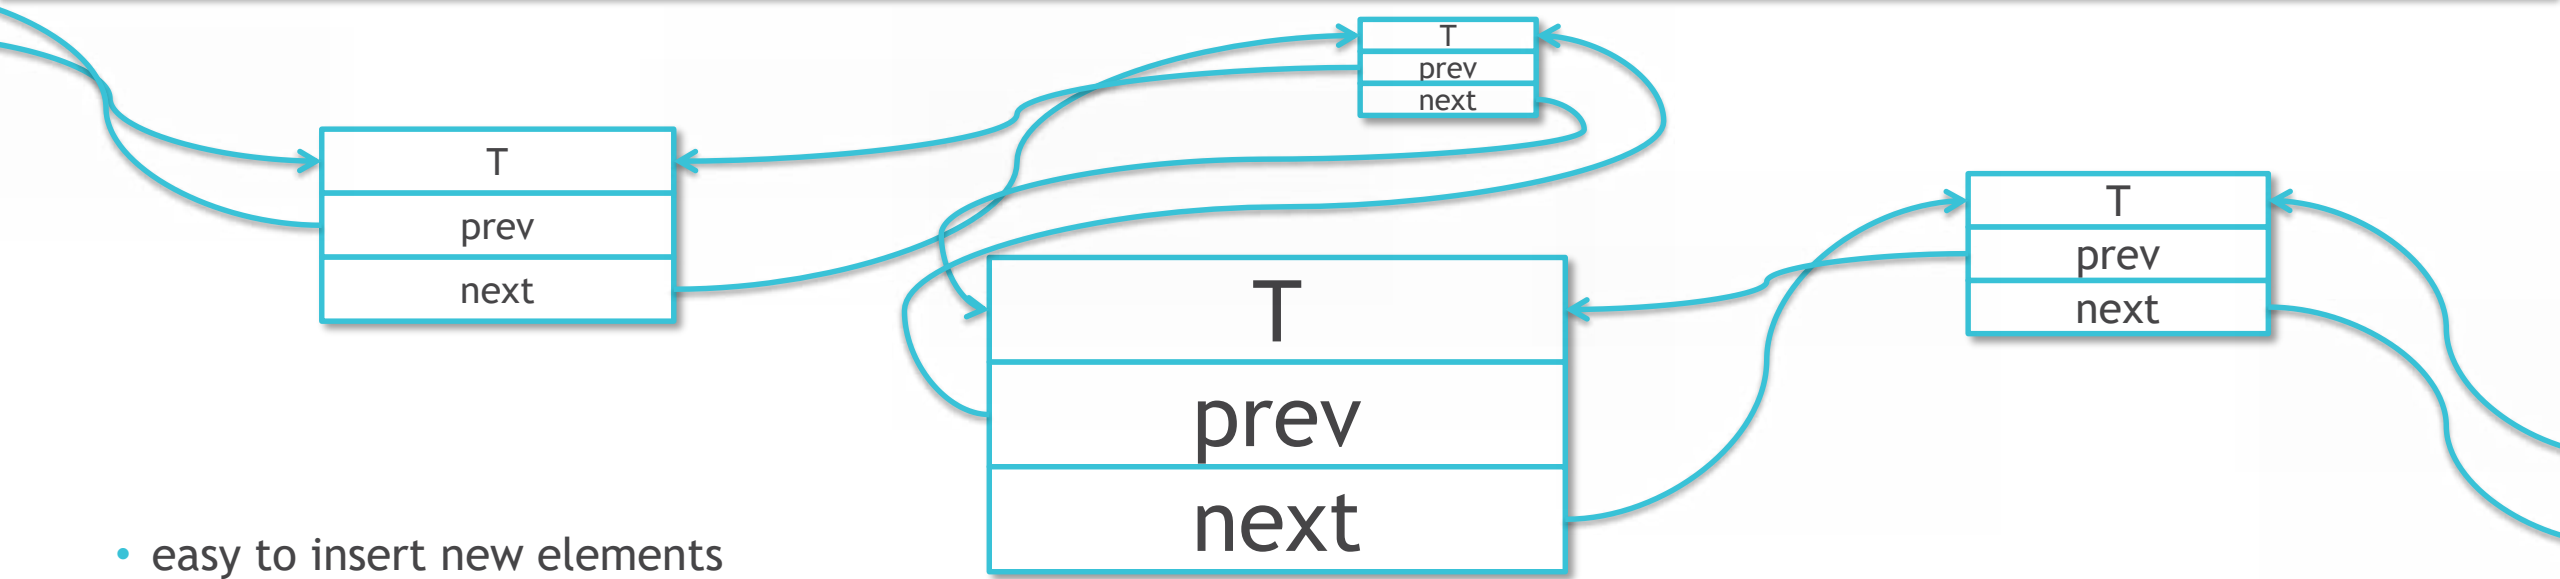
Operation	<code>std::list<T></code>	<code>std::vector<T></code>
create container of 'n' random values	$O(n)$	$O(n)$
sort container	$O(n \log n)$	$O(n \log n)$
'k' times insert a new random value keeping the container sorted	$O(k * (n+k))$	$O(k * (n+k))$
verify if sorted	$O(n+k)$	$O(n+k)$
Overall complexity	$O(k * (n+k))$	$O(k * (n+k))$

12x-44x speedup. Why?

std::list<T>




std::list<T>



- easy to insert new elements
- easy to erase existing elements
- easy to reorder elements
- bidirectional iteration
- memory ineffective
 - for small `T` large memory overhead
 - a lot of dynamic memory allocations and deallocations
 - a lot of pointer dereferences during iteration
 - not cache friendly

std::list<int> ITERATION

CPU cache



M
E
M
O
R
Y


std::vector<T>



- inserting new elements
 - end - easy (unless buffer reallocation is needed)
 - begin, middle - hard
- erasing existing elements
 - end - easy
 - begin, middle - hard
- swapping values needed to reorder elements
- random access iteration
- memory effective
 - small memory overhead (reserved space)
 - nearly no dynamic memory allocations and deallocations
 - no pointer dereferences during iteration
 - cache friendly!!!

std::vector<int> ITERATION

CPU cache



Writing cache friendly code makes a huge difference!

M
E
M
O
R
Y

LIST

Performance counter stats for './test_list':

83991,294231	task-clock (msec)	#	0,999 CPUs utilized
213901565783	cycles	#	2,547 GHz
27748493084	instructions	#	0,13 insn per cycle
8472128537	branches	#	100,869 M/sec
61763769	branch-misses	#	0,73% of all branches
8822212776	L1-dcache-loads	#	105,037 M/sec
4799991469	L1-dcache-load-misses	#	54,41% of all L1-dcache hits
4177763046	LLC-loads	#	49,740 M/sec
550506468	LLC-load-misses	#	13,18% of all LL-cache hits
78237	page-faults	#	0,931 K/sec

VECTOR

Performance counter stats for './test_vector':

2826,385943	task-clock (msec)	#	0,999 CPUs utilized
7462656329	cycles	#	2,640 GHz
27117855932	instructions	#	3,63 insn per cycle
7396349132	branches	#	2616,893 M/sec
39248559	branch-misses	#	0,53% of all branches
6466518567	L1-dcache-loads	#	2287,911 M/sec
77800147	L1-dcache-load-misses	#	1,20% of all L1-dcache hits
89304	LLC-loads	#	0,032 M/sec
36469	LLC-load-misses	#	40,84% of all LL-cache hits
1050	page-faults	#	0,371 K/sec

Latency Numbers Every Programmer Should Know

L1 cache reference	0.5	ns	
Branch misprediction	5	ns	
L2 cache reference	7	ns	14x L1 cache
Mutex lock/unlock	25	ns	
Main memory reference	100	ns	20x L2 cache, 200x L1 cache
Compress 1K bytes with Zippy	3,000	ns	
Send 1K bytes over 1 Gbps network	10,000	ns	
Read 4K randomly from SSD	150,000	ns	
Read 1 MB sequentially from memory	250,000	ns	
Round trip within same datacenter	500,000	ns	
Read 1 MB sequentially from SSD	1,000,000	ns	4X memory
Disk seek	10,000,000	ns	20x datacenter roundtrip
Read 1 MB sequentially from disk	20,000,000	ns	80x memory, 20X SSD
Send packet CA->Netherlands->CA	150,000,000	ns	

-- Jeff Dean

A NEW POSSIBILITY FOR ALGORITHMIC OPTIMIZATION

```
std::uint8_t test(int n, int k)
{
    // create collection of 'n' random samples
    // sort the collection
    // k times insert a new random sample keeping the collection sorted
    // return average value from samples in the collection
}
```

§25.4.3.1

template<class ForwardIterator, class T>

ForwardIterator lower_bound(ForwardIterator first, ForwardIterator last, const T& value);

- 1 Requires: The elements e of $[first, last)$ shall be partitioned with respect to the expression $e < value$*
- 2 Returns: The furthestmost iterator i in the range $[first, last]$ such that for every iterator j in the range $[first, i)$ the following corresponding conditions hold: $*j < value$.*
- 3 Complexity: At most $\log_2(last - first) + O(1)$ comparisons.*

CHANGE #3

Know and use C++ algorithms

```
std::uint8_t test(int n, int k)
{
    // ...
    // k times insert a new random sample keeping the collection sorted
    for(int i=0; i<k; ++i) {
        const sample new_sample{generate()};
        samples.emplace(lower_bound(begin(samples), end(samples), new_sample), new_sample);
    }
    // ...
}
```

1x-1.7x speedup

Whole code

Isn't it a great
example of efficient
programming?

```
collection init_random(int n, int k)
{
    collection samples;
    samples.reserve(n + k);
    for(int i=0; i<n; ++i)
        samples.emplace_back(generate());
    return samples;
}

std::uint8_t test(int n, int k)
{
    auto samples = init_random(n, k);
    sort(begin(samples), end(samples));

    // k times insert a new random sample keeping the collection sorted
    for(int i=0; i<k; ++i) {
        const sample new_sample{generate()};
        samples.emplace(lower_bound(begin(samples), end(samples), new_sample), new_sample);
    }

    return std::accumulate(begin(samples), end(samples), std::uint64_t{}) / samples.size();
}
```

QUOTES FROM THE C++ COMMUNITY

Using std::map is an exercise in slowing down code.
If you think link list is bad... maps are worse.

std::multiset<sample> usage

	test(10 000 000, 0)		test(0, 100 000)		test(100 000, 10 000)	
Metric	vector	multiset	vector	multiset	vector	multiset
task-clock [msec]	765	9162				
Instructions [insn/cycle]	1,33	0,31				
L1-dcache-loads [M/sec]	645,8	189,6				
L1-dcache-load-misses	0,39%	29,91%				
LLC-loads [M/sec]	0,074	37,9				
LLC-load-misses	37,55%	18,61%				
page-faults [K/sec]	1	11				

std::multiset<sample> usage

	test(10 000 000, 0)		test(0, 100 000)		test(100 000, 10 000)	
Metric	vector	multiset	vector	multiset	vector	multiset
task-clock [msec]	765	9162				
Instructions [insn/cycle]	1,33	0,31	3,86	0,60		
L1-dcache-loads [M/sec]	645,8	189,6	2557,5	326,3		
L1-dcache-load-misses	0,39%	29,91%	1,03%	12,06%		
LLC-loads [M/sec]	0,074	37,9	0,035	17,7		
LLC-load-misses	37,55%	18,61%	22,45%	16,84%		
page-faults [K/sec]	1	11	0,14	23		

std::multiset<sample> usage

	test(10 000 000, 0)		test(0, 100 000)		test(100 000, 10 000)	
Metric	vector	multiset	vector	multiset	vector	multiset
task-clock [msec]	765	9162	978	46		
Instructions [insn/cycle]	1,33	0,31	3,86	0,60		
L1-dcache-loads [M/sec]	645,8	189,6	2557,5	326,3		
L1-dcache-load-misses	0,39%	29,91%	1,03%	12,06%		
LLC-loads [M/sec]	0,074	37,9	0,035	17,7		
LLC-load-misses	37,55%	18,61%	22,45%	16,84%		
page-faults [K/sec]	1	11	0,14	23		

std::multiset<sample> usage

	test(10 000 000, 0)		test(0, 100 000)		test(100 000, 10 000)	
Metric	vector	multiset	vector	multiset	vector	multiset
task-clock [msec]	765	9162	978	46	212	61
Instructions [insn/cycle]	1,33	0,31	3,86	0,60	3,70	0,59
L1-dcache-loads [M/sec]	645,8	189,6	2557,5	326,3	2598,5	367,6
L1-dcache-load-misses	0,39%	29,91%	1,03%	12,06%	1,45%	19,20%
LLC-loads [M/sec]	0,074	37,9	0,035	17,7	0,084	16,9
LLC-load-misses	37,55%	18,61%	22,45%	16,84%	21,07%	8,38%
page-faults [K/sec]	1	11	0,14	23	0,67	23

TWEAKING 'N' IN 'F(N)'

Can we make our code faster?

Solution	test(10 000 000, 0)	test(0, 100 000)	test(100 000, 10 000)
Original	10,7s	104,5s	32,6s
Change #1	9,3s	57,1s	16,4s
Change #2	0,75s	1,7s	0,37s
Change #3	0,75s	0,98s	0,21s
Speedup	14x	106x	155x

TWEAKING 'N' IN 'F(N)'

Can we make our code faster?

Solution	test(10 000 000, 0)	test(0, 100 000)	test(100 000, 10 000)
Original	10,7s	104,5s	32,6s
Change #1	9,3s	57,1s	16,4s
Change #2	0,75s	1,7s	0,37s
Change #3	0,75s	0,98s	0,21s
Bonus #1	8,9s	0,04s	0,05s
Bonus #2	1,7s	0,04s	0,02s
Speedup	???	2 612x	1 630x

Bonus #1

std::multiset<sample> usage

```
using collection = std::multiset<elem_type>;  
// ...  
  
std::uint8_t test(int n, int k)  
{  
    // ...  
    sort(begin(samples), end(samples));  
    // ...  
    samples.emplace(lower_bound(begin(samples), end(samples), new_sample), new_sample);  
    // ...  
}
```

0.08x-24x speedup

Bonus #2

std::multiset<sample> + std::vector<sample> usage


```
using collection1 = std::vector<elem_type>;    using collection2 = std::multiset<elem_type>;  
// ...  
  
std::uint8_t test(int n, int k)  
{  
    // create collection of 'n' random samples  
    collection1 samples1 = init_random(n, k);  
  
    // sort the collection  
    sort(begin(samples1), end(samples1));  
  
    collection2 samples2{begin(samples1), end(samples1)};  
    // ...  
}
```

TAKE AWAYS



Learn C++ to be a more efficient programmer

- containers,
- algorithms,
- new C++11/C++14/C++17 features



Use RAII design pattern as
your first resort tool for
management of all
resources



Use tools provided by the
C++ standard or C++
community

Do not reinvent the wheel!!!



Remember that memory access is the bottleneck of many today's applications!


Limit pointers usage
Write cache friendly code




Photograph by Mauro Sartori

Always measure your code
if you care for performance!

Do not do premature
optimizations based only on
your assumptions
(algorithms, CPU cache
friendliness or others)





The background is a solid yellow color. It is decorated with several black geometric shapes, including triangles and parallelograms, arranged in a pattern that resembles a stylized border or a series of overlapping planes.

CAUTION
Programming
is addictive
(and too much fun)

A promotional image for the movie John Wick. Keanu Reeves is shown from the chest up, wearing a black leather jacket and dark sunglasses. He is holding a silver handgun in his right hand, pointing it towards the camera. The background is black with several bright red lines radiating outwards from behind him, creating a starburst effect. A light blue rectangular box is overlaid on the left side of the image, containing the text "BONUS SLIDE" in white capital letters.

BONUS SLIDE

PSEUDO-RANDOM NUMBERS GENERATOR

```
template<typename T>
auto make_random_int_generator(T min = std::numeric_limits<T>::min(),
                               T max = std::numeric_limits<T>::max(),
                               std::mt19937::result_type seed = std::mt19937::default_seed)
{
    static_assert(std::is_integral_v<T>);
    std::mt19937 gen{seed};
    std::uniform_int_distribution<T> distr{min, max};
    return [=]() mutable { return distr(gen); };
}

std::uint8_t generate()
{
    static auto generator = make_random_int_generator<std::uint8_t>();
    return generator();
}
```