


---

VIRTUALCONFERENCE // NOVEMBER2020 // PRESENTED BY EPAM





# High-level Abstractions, Safety, or Performance

**Mateusz Pusz**

---

VIRTUALCONFERENCE // NOVEMBER2020 // PRESENTED BY EPAM



# Who am I?

ADAPT BY ZED



## Mateusz Pusz

Associate Principal Software Engineer  
Head of the C++ Competency Center  
Head of EPAM Global C++ Community

### Modern C++ Evangelist

- Hacking C++ for more than 15 years for fun and living
- Trainer, coach, mentor, consultant, and conference speaker
- Active voting member and contributor of the ISO C++ Committee (WG21)
- MISRA C++ member
- Mainly interested in code performance, low latency, safety, and maintainability



**EPAM** ENGINEERING DNA



I am not claiming that C++ is the best programming language out there. Each of the programming languages has its use and areas of applicability where it proves the best. Let us not start the holy wars here 😊

mpusz Merge pull request #3 from MysterionRise/main · baafa5a · 3 hours ago · 18 commits

File	Commit Message	Time Ago
c++	Benchmarks updated	4 hours ago
java	Merge branch 'main' into main	3 hours ago
python	Benchmarks updated	4 hours ago
.clang-format	Initial version	yesterday
.gitattributes	.gitignore and .gitattributes added	2 days ago
.gitignore	clean up code, add simple jmh bench	21 hours ago
LICENSE	Initial commit	3 days ago
README.md	README updated	4 hours ago

**README.md**

## High-level Abstractions, Safety, or Performance?

This project provides a source code for my talk at EPAM's Zed 2020 conference.

During the talk I compare safety, user experience, and performance of a similar scenario in 3 different programming languages: Java, Python, and C++. The discussed feature is the physical units library.

### Libraries

- Java: JSR 385 - Units of Measurement
- Python: Pint
- C++: mp-units

**About**

Source code for EPAM Zed 2020 conference

Readme

MIT License

**Releases**

No releases published  
[Create a new release](#)

**Packages**

No packages published  
[Publish your first package](#)

**Contributors** 2

- mpusz Mateusz Pysz
- MysterionRise Konstantin Perikov

**Languages**


Language	Percentage
Java	35.3%
C++	33.2%
Python	25.1%
CMake	6.4%

- 
- 1** **Motivation**
  - 2** High-level abstractions
  - 3** Safety
  - 4** Efficiency




# The Mars Climate Orbiter

- Robotic space probe launched by NASA on December 11, 1998
- Project costs: \$327.6 million
  - spacecraft development: \$193.1 million
  - launching it: \$91.7 million
  - mission operations: \$42.8 million
- Mars Climate Orbiter began the planned orbital insertion maneuver on September 23, 1999 at 09:00:46 UTC








# What went wrong?

ADAPT BY ZED

- The **primary cause** of this discrepancy was that
  - one piece of ground software supplied by Lockheed Martin produced results in a **United States customary unit**, contrary to its Software Interface Specification (SIS)
  - second system, supplied by NASA, expected those results to be in **SI units**, in accordance with the SIS
- Specifically
  - software that calculated the total impulse produced by thruster firings calculated results in **pound-seconds**
  - the trajectory calculation software then used these results to update the predicted position of the spacecraft and expected it to be in **newton-seconds**





A long time ago in a galaxy far far away... ADAPT BY ZED





# Tactical Flight Computer

ADAPT BY ZED



# What is the correct order?

```
void DistanceBearing(double lat1, double lon1,  
                    double lat2, double lon2,  
                    double *Distance, double *Bearing);
```

```
void FindLatitudeLongitude(double Lat, double Lon,  
                          double Bearing, double Distance,  
                          double *lat_out, double *lon_out);
```

# double - an ultimate type

ADAPT BY ZED

```
double GlidePolar::MacCreadyAltitude(double emcready,  
                                     double Distance,  
                                     const double Bearing,  
                                     const double WindSpeed,  
                                     const double WindBearing,  
                                     double *BestCruiseTrack,  
                                     double *VMacCready,  
                                     const bool isFinalGlide,  
                                     double *TimeToGo,  
                                     const double AltitudeAboveTarget,  
                                     const double cruise_efficiency,  
                                     const double TaskAltDiff);
```



# Code we shouldn't write

```
// Air Density(kg/m3) from relative humidity(%),  
// temperature(°C) and absolute pressure(Pa)  
double AirDensity(double hr, double temp, double abs_press)  
{  
    return (1/(287.06 * (temp + 273.15))) *  
           (abs_press - 230.617 * hr * exp((17.5043 * temp)/(241.2 + temp)));  
}
```

# Now it is more important then ever

ADAPT BY ZED





# Comparison Scenario

ADAPT BY ZED

- Implement `avg_speed` function that takes `length` and `time` arguments and returns `speed` in the unit derived from the units of function arguments
- Calculate `avg_speed(220 km, 2 h)` and print the result in `km/h` and `m/s`
- Calculate `avg_speed(140 mi, 2 h)` and print the result in `mi/h` and `m/s`

- **Python** package to define, operate and manipulate physical quantities
- It allows **arithmetic operations** between them and **conversions from and to different units**
- It is distributed with a comprehensive **list of physical units, prefixes and constants**.



# Pint

```
ureg = UnitRegistry()

@ureg.check(['length'], ['time'])
def avg_speed(d, t):
    speed = d / t
    if not speed.check(['speed']):
        raise RuntimeError("Not a [speed] dimension")
    return speed


s1 = avg_speed(220 * ureg.kilometer, 2 * ureg.hour)
s2 = avg_speed(140 * ureg.mile, 2 * ureg.hour)

def print_si(q):
    print('{:~P}'.format(q))

print_si(s1)
print_si(s2)
print_si(s1.to('metre/second'))
print_si(s2.to_base_units())
```

```
110.0 km/hr
70.0 mi/hr
30.55555555555557 m/s
31.292800000000003 m/s
```

- The Unit of Measurement API
- Provides a set of **Java** language programming interfaces for handling units and quantities
  - checking of **unit compatibility**
  - **expression of a quantity** in various units
  - **arithmetic operations** on units






```
public static Quantity<Speed> avg_speed(Quantity<Length> length,
                                         Quantity<Time> time) throws ClassCastException {
    return length.divide(time).asType(Speed.class);
}
```

```
final Quantity<Speed> s1 = avg_speed(Quantities.getQuantity(220., KILO(Units.METRE)),
                                     Quantities.getQuantity(2., Units.HOUR));
final Quantity<Speed> s2 = avg_speed(Quantities.getQuantity(140., MILE),
                                     Quantities.getQuantity(2., Units.HOUR));

System.out.println(s1);
System.out.println(s2);
System.out.println(s1.to(Units.METRE_PER_SECOND));
System.out.println(s2.toSystemUnit());
```

- Modern C++ library
- Provides compile-time dimensional analysis and unit/quantity manipulation
  - no runtime execution or memory storage space cost is introduced
- Support for quantities and units for arbitrary unit system models and arbitrary value types
- Planned for the ISO standardization as a part of the C++23/26



```
constexpr Speed auto avg_speed(Length auto d, Time auto t)
{
    return d / t;
}

Speed auto s1 = avg_speed(si::length<si::kilometre>(220), si::time<si::hour>(2));
Speed auto s2 = avg_speed(si::length<si::international::mile>(140), si::time<si::hour>(2));
```

```
constexpr Speed auto avg_speed(Length auto d, Time auto t)
{
    return d / t;
}
```

```
Speed auto s1 = avg_speed(220_q_km, 2_q_h);
Speed auto s2 = avg_speed(140_q_mi, 2_q_h);
```

```
std::cout << s1 << '\n';
std::cout << s2 << '\n';
std::cout << quantity_cast<si::metre_per_second>(s1) << '\n';
std::cout << quantity_cast<si::dim_speed::coherent_unit>(s2) << '\n';
```

```
110 km/h
70 mi/h
30.5556 m/s
31.2928 m/s
```

- 1 Motivation
- 2 High-level abstractions
- 3 Safety**
- 4 Efficiency

- Ensure that for reordered arguments `avg_speed(2 h, 220 km)` returns an error
- Ensure that an error is reported when `avg_speed` returns the result of an invalid calculation
  - function multiplies the arguments instead of dividing them
  - the result is not a quantity of speed

# Pint – reordered arguments

```
@ureg.check('[length]', '[time]')
def avg_speed(d, t):
    speed = d / t
    if not speed.check('[speed]'):
        raise RuntimeError("Not a [speed] dimension")
    return speed

s1 = avg_speed(2 * ureg.hour, 220 * ureg.kilometer)
```

Traceback (most recent call last):

```
File "safety_1.py", line 13, in <module>
    s1 = avg_speed(2 * ureg.hour, 220 * ureg.kilometer)
File "/home/mpusz/.local/lib/python3.8/site-packages/pint/registry_helpers.py", line 350, in wrapper
    raise DimensionalityError(value, "a quantity of", val_dim, dim)
pint.errors.DimensionalityError: Cannot convert from '2 hour' ([time]) to 'a quantity of' ([length])
```

Runtime Error



# Pint – invalid computation

ADAPT BY ZED

```
@ureg.check('[length]', '[time]')
def avg_speed(d, t):
    speed = d * t
    if not speed.check('[speed]'):
        raise RuntimeError("Not a [speed] dimension")
    return speed

s1 = avg_speed(220 * ureg.kilometer, 2 * ureg.hour)
```

---

```
Traceback (most recent call last):
  File "safety_2.py", line 14, in <module>
    s1 = avg_speed(220 * ureg.kilometer, 2 * ureg.hour)
  File "/home/mpusz/.local/lib/python3.8/site-packages/pint/registry_helpers.py", line 351, in wrapper
    return func(*args, **kwargs)
  File "safety_2.py", line 10, in avg_speed
    raise RuntimeError("Not a [speed] dimension")
RuntimeError: Not a [speed] dimension
```

Runtime Error

# JSR 385 – reordered arguments

ADAPT BY ZED

```
public static Quantity<Speed> avg_speed(Quantity<Length> length,
                                         Quantity<Time> time) throws ClassCastException {
    return length.divide(time).asType(Speed.class);
}

public static void main(String[] args) {
    final Quantity<Speed> s = avg_speed(Quantities.getQuantity(2., Units.HOUR),
                                         Quantities.getQuantity(220., KILO(Units.METRE)));
}
```

---

[ERROR] Safety\_1.java:[41,31] incompatible types: inference variable Q has incompatible equality constraints  
javax.measure.quantity.Length, javax.measure.quantity.Time

Compile-time Error

# JSR 385 – invalid computation

ADAPT BY ZED

```
public static Quantity<Speed> avg_speed(Quantity<Length> length,
                                         Quantity<Time> time) throws ClassCastException {
    return length.multiply(time).asType(Speed.class);
}

public static void main(String[] args) {
    final Quantity<Speed> s = avg_speed(Quantities.getQuantity(220., KILO(Units.METRE)),
                                         Quantities.getQuantity(2., Units.HOUR));
}
```

---

```
Exception in thread "main" java.lang.ClassCastException: The unit: km·h is not compatible with quantities of type
interface javax.measure.quantity.Speed
    at tech.units.indriya.AbstractUnit.asType(AbstractUnit.java:277)
    at tech.units.indriya.AbstractUnit.asType(AbstractUnit.java:89)
    at tech.units.indriya.AbstractQuantity.asType(AbstractQuantity.java:337)
    at tech.units.indriya.AbstractQuantity.asType(AbstractQuantity.java:114)
    at zed2020.Safety_2.avg_speed(Safety_2.java:37)
    at zed2020.Safety_2.main(Safety_2.java:41)
```

Runtime Error

# mp-units – reordered arguments

ADAPT BY ZED

```
constexpr Speed auto avg_speed(Length auto d, Time auto t)
{
    return d / t;
}
```

```
Speed auto s1 = avg_speed(2_q_h, 220_q_km);
```

---

```
avg_speed.cpp: In function 'int main()':
avg_speed.cpp:45:44: error: use of function 'constexpr Speed auto avg_speed(auto:16, auto:17) [with auto:16 =
units::quantity<units::physical::si::dim_time, units::physical::si::hour, long int>; auto:17 =
units::quantity<units::physical::si::dim_length, units::physical::si::kilometre, long int>]' with unsatisfied
constraints
```

```
45 |     Speed auto s1 = avg_speed(2_q_h, 220_q_km);
    |                               ^
```

```
avg_speed.cpp:29:16: note: declared here
```

```
29 |     constexpr Speed auto avg_speed(Length auto d, Time auto t)
    |                               ^~~~~~
```

Compile-time Error

# mp-units – invalid computation

ADAPT BY ZED

```
constexpr Speed auto avg_speed(Length auto d, Time auto t)
{
    return d * t;
}
```

```
Speed auto s1 = avg_speed(220_q_km, 2_q_h);
```

---

In instantiation of 'constexpr auto [requires units::physical::Speed<<placeholder>, >] avg\_speed(auto:16, auto:17) [with auto:16 = units::quantity<units::physical::si::dim\_length, units::physical::si::kilometre, long int>; auto:17 = units::quantity<units::physical::si::dim\_time, units::physical::si::hour, long int>]':

avg\_speed.cpp:42:44: required from here

avg\_speed.cpp:32:14: **error**: deduced return type does not satisfy placeholder constraints

```
32 |     return d * t;
    |           ^
```

Compile-time Error

- 1 Motivation
- 2 High-level abstractions
- 3 Safety
- 4 **Efficiency**

- Benchmark the following scenarios both for operations on fundamental/primitive types and on high-level quantity abstractions
  - **Arithmetic** - create quantities of **length** and **time** and divide them to obtain **speed**
  - **Scaling** - create a quantity of **speed** and convert the unit from **km/h** to **m/s**




Pint can impose a significant performance overhead on computationally-intensive problems. The following are some suggestions for getting the best performance.


## Use magnitudes when possible

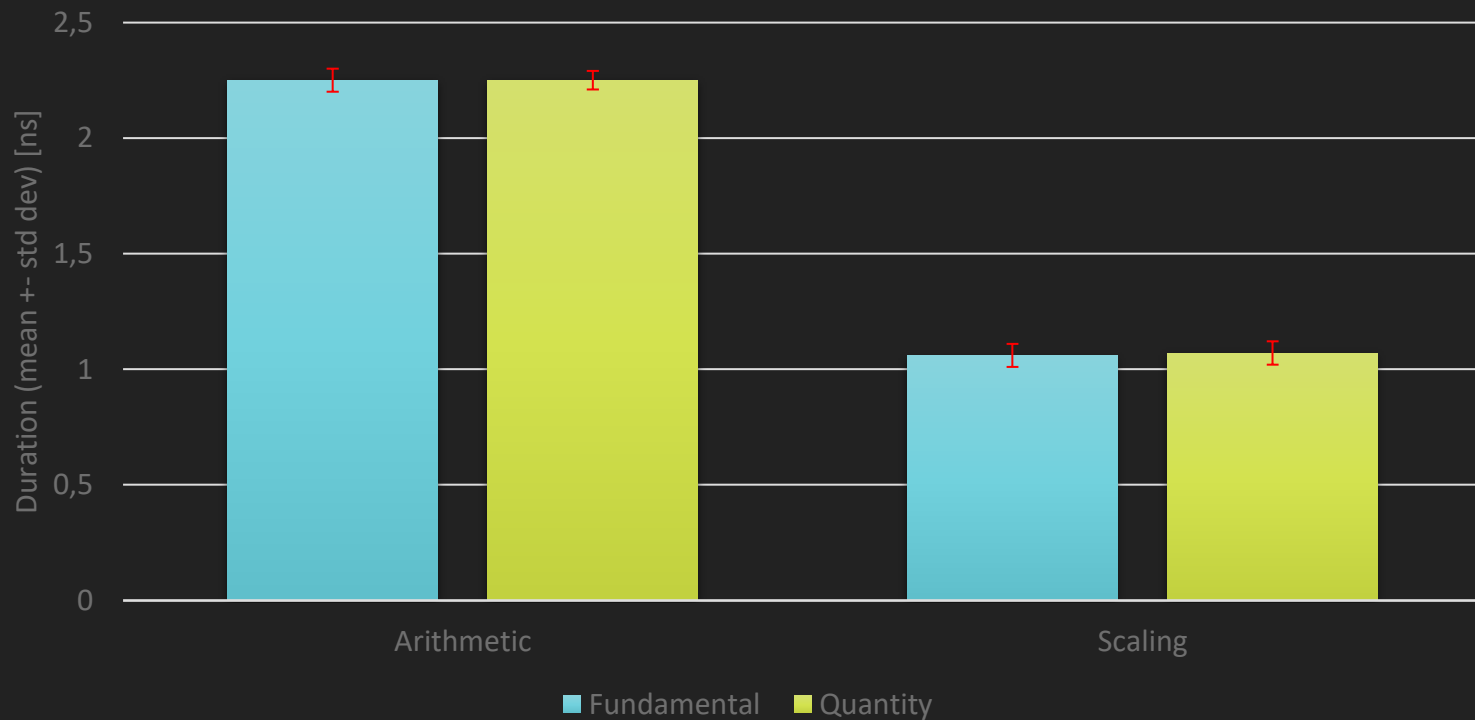
It's significantly faster to perform mathematical operations on magnitudes (even though you're still using pint to retrieve them from a quantity object).

```
In [1]: from pint import UnitRegistry  
  
In [2]: ureg = UnitRegistry()  
  
In [3]: q1 = ureg('1m')  
  
In [5]: q2 = ureg('2m')  
  
In [6]: %timeit (q1-q2)  
10000 loops, best of 3: 7.9 µs per loop  
  
In [7]: %timeit (q1.magnitude-q2.magnitude)  
100000 loops, best of 3: 356 ns per loop
```

Bear in mind that altering computations like this **loses the benefits of automatic unit conversion**, so use with care.







```
constexpr double avg_speed(double d, double t)
{
    return d / t;
}
```

```
constexpr Speed auto avg_speed(Length auto d, Time auto t)
{
    return d / t;
}
```

```
avg_speed(...):
    divsd    xmm0, xmm1
    ret
```

```
avg_speed(...):
    divsd    xmm0, xmm1
    ret
```

High-level abstractions without sacrificing runtime performance

# Compile-time evaluation when possible

ADAPT BY ZED

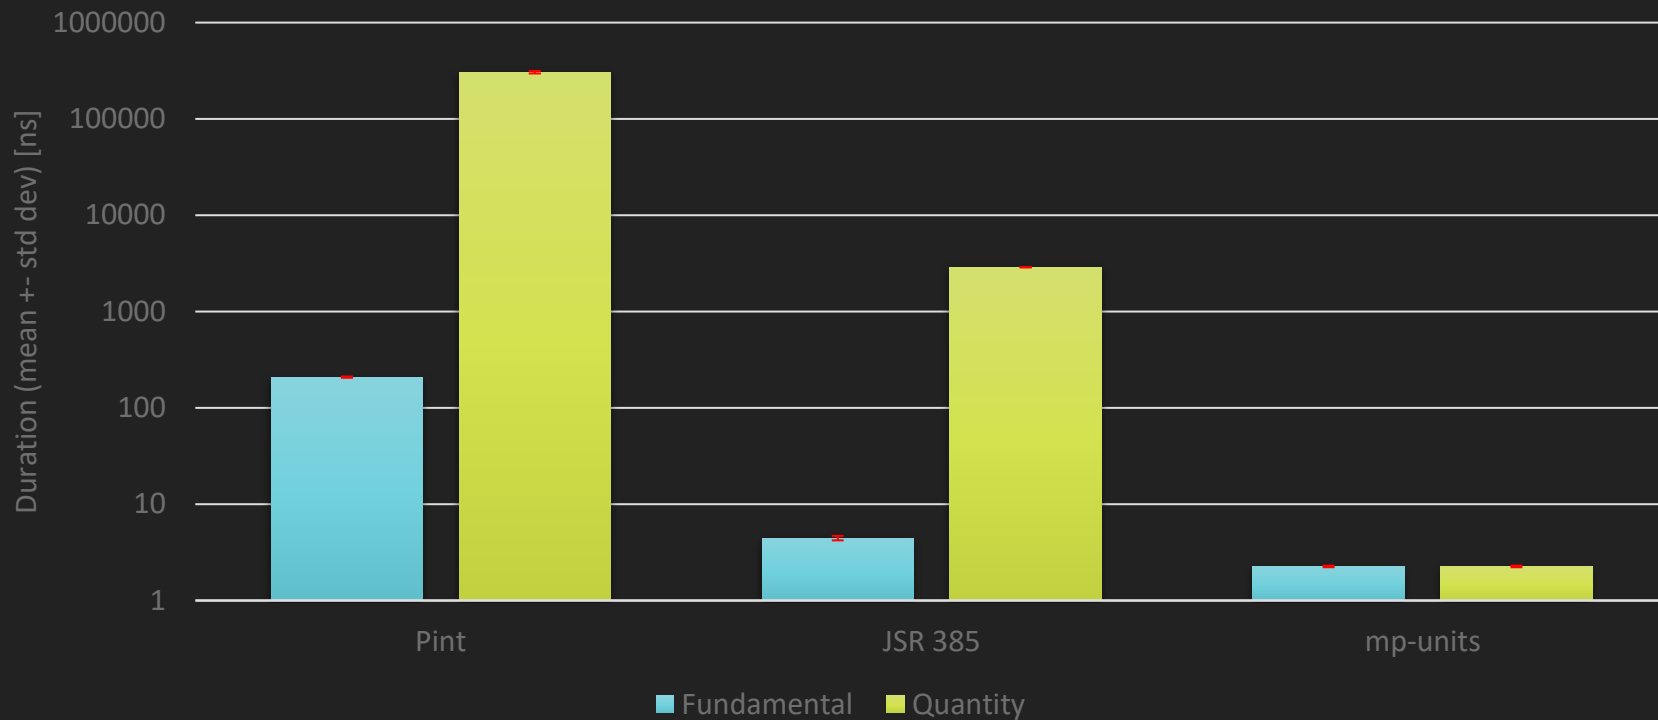
```
// simple numeric operations
static_assert(10_q_km / 2 == 5_q_km);

// unit conversions
static_assert(1_q_h == 3600_q_s);
static_assert(1_q_km + 1_q_m == 1001_q_m);

// dimension conversions
static_assert(1_q_km / 1_q_s == 1000_q_m_per_s);
static_assert(2_q_km_per_h * 2_q_h == 4_q_km);
static_assert(2_q_km / 2_q_km_per_h == 1_q_h);
static_assert(2_q_m * 3_q_m == 6_q_m2);
static_assert(10_q_km / 5_q_km == 2);
static_assert(1000 / 1_q_s == 1_q_kHz);
```

# Comparison (Arithmetic)

ADAPT BY ZED





Efficiency is not just running fast or running bigger programs, it's also running using less resources.

Bjarne Stroustrup, June 2011



Quantity
<code>_magnitude</code>
<code>_units</code>
<code>__used</code>
<code>__handling</code>

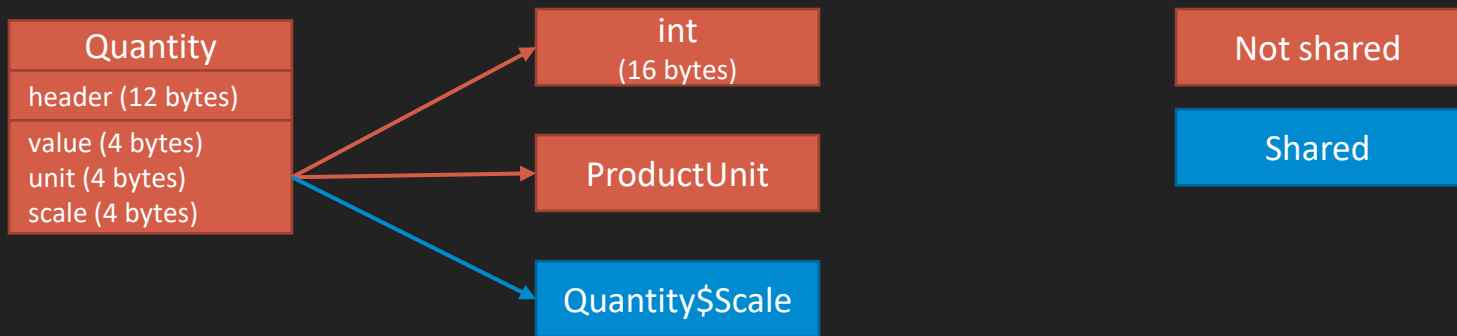
```
from pymboler import asizeof

def test_data(count, min, max):
    values = default_rng().uniform(min, max, count)
    return list(map(lambda s: s * ureg.kilometre / ureg.hour, values))

print(asizeof.asizeof(test_data(10000, 40, 140)), "bytes")
```

6963520 bytes

1 Quantity ~ 700 bytes



```
Quantity[] quantities = new Quantity[10000];
for (int i = 0; i < quantities.length; ++i) {
    quantities[i] = avg_speed(Quantities.getQuantity(Math.random() * 1000, KILO(Units.METRE)),
                             Quantities.getQuantity(Math.random() * 10, Units.HOUR));
}
System.out.println(RamUsageEstimator.sizeOf(quantities) + " bytes");
```

5001128 bytes

1 Quantity ~ 500 bytes


```
quantity<Dimension D, UnitOf<D> U, QuantityValue Rep>
```

Rep value

```
namespace si {  
  
template<UnitOf<dim_length> U, QuantityValue Rep = double>  
using length = quantity<dim_length, U, Rep>;  
  
}  
  
static_assert(sizeof(si::length<si::metre>) == sizeof(double));  
static_assert(sizeof(si::length<si::metre, int>) == sizeof(int));  
static_assert(sizeof(si::length<si::metre, std::int8_t>) == 1);
```



Modern C++ provides safety and  
high-level abstractions without  
sacrificing on efficiency.



# Thank you!

---

VIRTUALCONFERENCE // NOVEMBER2020 // PRESENTED BY EPAM

